

# CS585 Problem Set 4 (Total points: 35 + 5 bonus)

Assignment adapted from Svetlana Lazebnik

## Instructions

1. Assignment is due at **5 PM on Tuesday Mar 29 2022**.

2. Submission instructions:

A. A single `.pdf` report that contains your work for Q1, Q2, Q3, and Q4. For Q1 you can either type out your responses in LaTeX, or any other word processing software. You can also hand write them on a tablet, or scan in hand-written answers. If you hand-write, please make sure they are neat and legible. If you are scanning, make sure that the scans are legible. Lastly, convert your work into a `PDF`.

For Q2, Q3, and Q4 your response should be electronic (no handwritten responses allowed). You should respond to the questions Q2, Q3, and Q4 individually and include images as necessary. Your response to all questions in the PDF report should be self-contained. It should include all the output you want us to look at. You will not receive credit for any results you have obtained, but failed to include directly in the PDF report file.

PDF file should be submitted to [Gradescope](#) under `PS4`. Please tag the responses in your PDF with the Gradescope questions outline as described in [Submitting an Assignment](#).

B. You also need to submit code for Q2, Q3, and Q4 in the form of a single `.zip` file that includes all your code, all in the same directory. You can submit Python code in `.ipynb` format. Code should also be submitted to [Gradescope](#) under `PS4-Code`. *Not submitting your code will lead to a loss of 100% of the points on Q2, Q3, and Q4.*

C. We reserve the right to take off points for not following submission instructions. In particular, please tag the responses in your PDF with the Gradescope questions outline as described in [Submitting an Assignment](#).

## Problems

### 1. Optical Flow [10 pts, 3 parts]

A. [2 pts] Describe a scenario where the object is not moving but the optical flow field is not zero.

If the lighting changes, the optical flow field won't be zero even though the object is stationary.

A. [3 pts] The Constant Brightness Assumption (CBA) is used in the Lucas and Kanade Algorithm. Describe how the algorithm could be adjusted to handle the fact that the

assumption might be violated.

The algorithm can use normalized correlation when matching feature patches to relax the brightness constancy requirement.

A. **[5 pts]** Why does the first order Taylor series provide a reasonable approximation for estimating optical flow?

We assume that changes in motion will be small from frame to frame, so a first order Taylor series is sufficient. We use the Taylor expansion because it can linearize the equations we want to solve for.

1. **Camera Calibration [10 pts]**. For the pair of images in the folder `calibraion`, calculate the camera projection matrices by using 2D matches in both views and 3D point coordinates in `lab_3d.txt`. Once you have computed your projection matrices, you can evaluate them using the provided evaluation function (`evaluate_points`). The function outputs the projected 2-D points and residual error. Report the estimated  $3 \times 4$  camera projection matrices (for each image), and residual error. **Hint:** The residual error should be  $< 20$  and the squared distance of the projected 2D points from actual 2D points should be  $< 4$ .

In [1]:

```
from PIL import Image
import numpy as np
import matplotlib.pyplot as plt

def evaluate_points(M, points_2d, points_3d):
    """
    Visualize the actual 2D points and the projected 2D points calculated from
    the projection matrix
    You do not need to modify anything in this function, although you can if you
    want to
    :param M: projection matrix 3 x 4
    :param points_2d: 2D points N x 2
    :param points_3d: 3D points N x 3
    :return:
    """
    N = len(points_3d)
    points_3d = np.hstack((points_3d, np.ones((N, 1))))
    points_3d_proj = np.dot(M, points_3d.T).T
    u = points_3d_proj[:, 0] / points_3d_proj[:, 2]
    v = points_3d_proj[:, 1] / points_3d_proj[:, 2]
    residual = np.sum(np.hypot(u-points_2d[:, 0], v-points_2d[:, 1]))
    points_3d_proj = np.hstack((u[:, np.newaxis], v[:, np.newaxis]))
    return points_3d_proj, residual

# Write your code here for camera calibration
def camera_calibration(pts_2d, pts_3d):
    """
    write your code to compute camera matrix
    """
    # <YOUR CODE>
    A = []
    for i in range(len(pts_3d)):
        u, v = pts_2d[i][0], pts_2d[i][1]
        x, y, z = pts_3d[i][0], pts_3d[i][1], pts_3d[i][2]
        row1 = [x, y, z, 1, 0, 0, 0, 0, -u*x, -u*y, -u*z, -u]
```

```

        row2 = [0, 0, 0, 0, x, y, z, 1, -v*x, -v*y, -v*z, -v]
        A.append(row1)
        A.append(row2)
    A = np.array(A)
    eig_vals, eig_vecs = np.linalg.eig(A.T @ A)
    p = eig_vecs[:, np.argmin(eig_vals)]
    P = np.vstack((p[:4], p[4:8], p[8:]))

    return P

# Load 3D points, and their corresponding locations in
# the two images.
pts_3d = np.loadtxt('calibration/lab_3d.txt')
matches = np.loadtxt('calibration/lab_matches.txt')

# print lab camera projection matrices:
lab1_proj = camera_calibration(matches[:, :2], pts_3d)
lab2_proj = camera_calibration(matches[:, 2:], pts_3d)
print('lab 1 camera projection')
print(lab1_proj)

print('')
print('lab 2 camera projection')
print(lab2_proj)

# evaluate the residuals for both estimated cameras
_, lab1_res = evaluate_points(lab1_proj, matches[:, :2], pts_3d)
print('residuals between the observed 2D points and the projected 3D points:')
print('residual in lab1:', lab1_res)
_, lab2_res = evaluate_points(lab2_proj, matches[:, 2:], pts_3d)
print('residual in lab2:', lab2_res)

```

```

lab 1 camera projection
[[ 3.09963781e-03  1.46206985e-04 -4.48502644e-04 -9.78930624e-01]
 [ 3.07018392e-04  6.37194274e-04 -2.77356105e-03 -2.04144664e-01]
 [ 1.67933458e-06  2.74767796e-06 -6.83970171e-07 -1.32882925e-03]]

```

```

lab 2 camera projection
[[ 6.93154111e-03 -4.01683849e-03 -1.32603329e-03 -8.26700612e-01]
 [ 1.54768631e-03  1.02452783e-03 -7.27440248e-03 -5.62523170e-01]
 [ 7.60945430e-06  3.70953989e-06 -1.90203701e-06 -3.38807512e-03]]

```

```

residuals between the observed 2D points and the projected 3D points:
residual in lab1: 13.545700151863898
residual in lab2: 15.544918932028565

```

1. **Camera Centers [5 pts]**. Calculate the camera centers using the estimated or provided projection matrices. Report the 3D locations of both cameras in your report. **Hint:** Recall that the camera center is given by the null space of the camera matrix.

In [2]:

```

import scipy.linalg
# Write your code here for computing camera centers
def calc_camera_center(proj):
    """
    write your code to get camera center in the world
    from the projection matrix
    """

```

```
# <YOUR CODE>
return scipy.linalg.null_space(proj)
```

```
# <YOUR CODE> compute the camera centers using
# the projection matrices
lab1_c = calc_camera_center(lab1_proj)
lab2_c = calc_camera_center(lab2_proj)
print('lab1 camera center', lab1_c)
print('lab2 camera center', lab2_c)
lab1_c = lab1_c.T/lab1_c[-1]
lab2_c = lab2_c.T/lab2_c[-1]
```

```
lab1 camera center [[0.7072697 ]
 [0.70349616]
 [0.06969489]
 [0.0023126  ]]
lab2 camera center [[0.70061827]
 [0.71005902]
 [0.07031994]
 [0.00231151]]
```

1. **Triangulation [10 pts].** Use linear least squares to triangulate the 3D position of each matching pair of 2D points using the two camera projection matrices. As a sanity check, your triangulated 3D points for the lab pair should match very closely the originally provided 3D points in `lab_3d.txt`. Display the two camera centers and reconstructed points in 3D. Include snapshots of this visualization in your report. Also report the residuals between the observed 2D points and the projected 3D points in the two images. Note: You do not need the camera centers to solve the triangulation problem. They are used just for the visualization.

In [3]:

```
# Write your code here for triangulation
from mpl_toolkits.mplot3d import Axes3D
def triangulation(lab_pt1, lab1_proj, lab_pt2, lab2_proj):
    """
    write your code to triangulate the points in 3D
    """
    # <YOUR CODE>
    reconst_points = []
    for i in range(len(lab_pt1)):
        x, y = lab_pt1[i][0], lab_pt1[i][1]
        x_prime, y_prime = lab_pt2[i][0], lab_pt2[i][1]
        p1, p2 = lab1_proj, lab2_proj
        A = np.vstack((y * p1[2] - p1[1], p1[0] - x * p1[2],
                       y_prime * p2[2] - p2[1], p2[0] - x_prime * p2[2]))
        eig_vals, eig_vecs = np.linalg.eig(A.T @ A)
        v = eig_vecs[:, np.argmin(eig_vals)]
        v = v / v[3]
        reconst_points.append(v[:3])
    return np.array(reconst_points)

def evaluate_points_3d(points_3d_lab, points_3d_gt):
    """
    write your code to evaluate the triangulated 3D points
    """
    # <YOUR CODE>
```

```

return np.sum((points_3d_lab - points_3d_gt) ** 2)

# triangulate the 3D point cloud for the lab data
matches_lab = np.loadtxt('calibration/lab_matches.txt')
lab_pt1 = matches_lab[:, :2]
lab_pt2 = matches_lab[:, 2:]
points_3d_gt = np.loadtxt('calibration/lab_3d.txt')
points_3d_lab = triangulation(lab_pt1, lab1_proj, lab_pt2, lab2_proj)
res_3d_lab = evaluate_points_3d(points_3d_lab, points_3d_gt)
print('Mean 3D reconstruction error for the lab data: ', round(np.mean(res_3d_lab), 5))
_, res_2d_lab1 = evaluate_points(lab1_proj, lab_pt1, points_3d_lab)
_, res_2d_lab2 = evaluate_points(lab2_proj, lab_pt2, points_3d_lab)
print('2D reprojection error for the lab 1 data: ', np.mean(res_2d_lab1))
print('2D reprojection error for the lab 2 data: ', np.mean(res_2d_lab2))
# visualization of lab point cloud
camera_centers = np.vstack((lab1_c, lab2_c))
fig = plt.figure()
ax = fig.add_subplot(111, projection='3d')
ax.scatter(points_3d_lab[:, 0], points_3d_lab[:, 1], points_3d_lab[:, 2], c='b', label=
ax.scatter(camera_centers[:, 0], camera_centers[:, 1], camera_centers[:, 2], c='g', s=5
ax.legend(loc='best')

```


Mean 3D reconstruction error for the lab data: 0.00492

2D reprojection error for the lab 1 data: 10.899312973796714

2D reprojection error for the lab 2 data: 1.5485007284731922

<matplotlib.legend.Legend at 0x2424166f6d0>

Out[3]:



1. **Extra Credits [5 pts].** Use the putative match generation and RANSAC code from PS3 to estimate fundamental matrices without ground-truth matches. For this part, only use the normalized algorithm. Report the number of inliers and the average residual for the inliers. Compare the quality of the result with the one you get from ground-truth matches.

In [ ]: